

Skills & Subagents Catalog Delta — June 2026

Anthropic · VoltAgent · shanraishan — what changed since 2026-05-28, with adoption verdicts and dynamic-workflow recommendations

Skills & Subagents Catalog Delta — June 2026

Window reviewed: 2026-05-28 (last §7 refresh) → 2026-06-19. **Sources:** Anthropic (anthropics/skills + anthropics/claude-plugins-official + Claude Code docs) · VoltAgent (awesome-claude-code-subagents) · shanraishan (claude-code-best-practice). **Method:** read-only research-scout fan-out + raw-file reads of every candidate agent. Confidence tags: [verified] = read directly this session; [inferred] = derived from listings/commits; [speculative] = unconfirmed.

1. Executive summary

The headline isn't a new agent — it's a **platform shift**. **Dynamic Workflows went GA (~2026-05-28)**: script-orchestrated fan-out of up to 16 concurrent / 1000 total subagents, results kept out of the context window, resumable, and saveable as slash commands. That reframes "better prompting for dynamic workflows" from a technique into operating a real orchestration runtime — captured as a new kit asset (rules/dynamic-workflow-prompting.md) and a new §7 orchestration tier.

Three things to act on:

- 1. **Four clean VoltAgent agents are worth adopting** — mcp-developer , cli-developer , first-principles-thinking , api-documenter . All hardened and added to the kit this pass. The rest of the ~23 new agents are team/infra/enterprise mismatches or carry the same context-manager-handshake / fabricated-metric anti-patterns that got most of the May batch skipped.
- 2. **shanraishan added genuinely useful prompting craft** — !command injection, PROACTIVELY triggers, Agent(agent_type) tool restriction, and an agent-teams data-contract pattern. Folded into the new asset.
- 3. **One flag for your decision:** content-quality-editor is the cleanest new agent structurally, but it does the exact job (unslap -based AI-tell stripping) that your global CLAUDE.md deprecated ai-writing-auditor for. Left unadopted pending your call.

Anthropic's skills repo is essentially unchanged (still 17); the action there is the platform features, not new installables.

2. Per-source deltas

2a. Anthropic (T1) — platform-level

Item	What changed	Confidence	Relevance
Dynamic Workflows GA	New first-class primitive (~05-28). ≤16 concurrent / 1000 total subagents; results in script vars; resumable; saved → slash command. Trigger: ultracode / /effort ultracode / /deep-research / "use a workflow" (v2.1.160+). Needs v2.1.154+.	verified	High — already in your environment; drives the prompting asset + §7 tier.
/deep-research	First bundled workflow command — adversarial multi-source research → cited report.	verified	High — candidate replacement for ad-hoc research fan-outs.

Item	What changed	Confidence	Relevance
Table 5 / Mythos 5	New model family (06-09), tuned for long multi-step autonomous tasks. Fable 5 (<code>claude-fable-5</code>) public; Mythos 5 restricted.	verified	High — preferred <code>model:</code> for hard workflow-orchestration stages.
claude-api skill refresh	Updated 06-09 with Fable 5 / Mythos 5 IDs, scheduled deployments, vault env-var credentials.	verified	Medium.
frontend-design skill refresh	Updated 06-07.	verified	Low.
project-artifact plugin	New (06-19): living HTML status page for multi-workstream projects → <code>claude.ai/code/artifact/...</code> . Needs the Artifact tool (Team/Enterprise beta).	verified	Situational.
anthropics/skills count	Still 17; none added/removed since baseline.	verified	— confirms no install action needed there.

Sources: `code.claude.com/docs/en/workflows` · `claude.com/blog/introducing-dynamic-workflows-in-claude-code` · `anthropic.com/news/claude-fable-5-mythos-5` · `github.com/anthropics/skills/commits/main` · `github.com/anthropics/claude-plugins-official/commits/main` .

2b. VoltAgent (T2) — catalog grew ~131 → ~154 agents

New agents cluster in **04-quality-security** (~8: accessibility-tester, ad-security-reviewer, chaos-engineer, gdpr-ccpa-compliance, penetration-tester, powershell-security-hardening, qa-expert, ui-ux-tester), **06-developer-experience** (~6: build-engineer, cli-developer, dependency-manager, **mcp-developer**, powershell-module-architect, powershell-ui-architect, visual-asset-generator), and **08-business-product** (assumption-mapping, backlog-grooming, content-quality-editor, growth-loops). [verified by directory listing]

No structural improvement: the fabricated-performance-metrics flavor text and context-manager JSON handshakes that flagged the May batch are still present in many new agents. [verified by raw-file read] A 06-15 commit removed VoltAgent branding from the README — possible governance change. [inferred]

Sources: repo README + per-category trees + raw files (see adoption matrix §3 for the files actually read).

2c. shanraishan (T2/T3) — richest for the dynamic-workflow ask

Item	What it is	Confidence
agent-teams/ directory <i>(new)</i>	Concrete pattern: data-contract negotiation between parallel teammates, parallel-component independence, deterministic output routing, frontmatter resource caps (<code>model: haiku</code> , <code>maxTurns: 3</code>).	verified
Five named macro-workflows	Superpowers, Everything Claude Code, Matt Pocock Skills, Spec Kit, gstack — all converge on Research → Plan → Execute → Review → Ship .	verified
!command injection	<code>!</code> -prefixed line runs a shell command and injects output into the prompt — feed live state (git diff, build error) instead of describing it.	inferred (in README; matches real CC feature)
PROACTIVELY trigger	The keyword in a subagent description biases auto-invocation.	inferred

Item	What it is	Confidence
Agent (agent_type) restriction	In the <code>tools:</code> field, restricts which subagents may be spawned.	inferred
"Hot Features" list + star counts	Lists ~20 features (<code>/sandbox</code> , <code>/focus</code> , <code>/advisor</code> , Auto Mode, etc.); some look aspirational and star counts look inflated.	speculative — do not cite uncritically.

Repo is actively maintained but mostly via automated daily chore commits; human content additions aren't cleanly datable. Sources: repo README + `agent-teams/agent-teams-prompt.md` + `best-practice/claude-subagents.md` + repo tree.

3. Adoption verdict matrix (VoltAgent new agents)

Judged against the kit's standing rules: explicit minimal `tools:` , no context-manager JSON handshakes, no fabricated-metric templates, read-only reviewers. All rows `[verified from raw file]` unless noted.

Agent	Tools	Anti-patterns present?	Verdict	Reason
mcp-developer	Read,Write,Edit,Bash,Glob,Grep	context-mgr opener + handshake + fake metrics (stripped)	ADOPTED	Only purpose-built MCP agent; fills May gap.
cli-developer	Read,Write,Edit,Bash,Glob,Grep	handshake + fake metrics (stripped)	ADOPTED	CLI design for kit/UE5/portfolio tooling.
first-principles-thinking	Read,Grep,Glob,WebFetch,WebSearch (read-only)	only cross-agent prose (stripped)	ADOPTED	Clean reasoning agent; no overlap with research-scout.
api-documenter	Read,Write,Edit,Glob,Grep,WebFetch,WebSearch	handshake + fake metrics (stripped); Write now scoped to docs	ADOPTED	OpenAPI + examples; complements technical-writer.
content-quality-editor	Read,Write,Edit,Bash	none structurally	FLAG — your call	Does the <code>unslop</code> AI-tell-stripping job your <code>CLAUDE.md</code> deprecated <code>ai-writing-auditor</code> for; <code>unslop</code> not installed.
accessibility-tester	Read,Grep,Glob,Bash (read-only)	likely context-mgr block (line ~95, unconfirmed)	DEFER	Confirm the block; adopt only if shipping a11y-sensitive web.
penetration-tester	Read,Grep,Glob,Bash (read-only, Opus)	context-mgr opener; rest unread	DEFER	Portfolio pre-launch security

Agent	Tools	Anti-patterns present?	Verdict	Reason
				sweep.
llm-architect	Read,Write,Edit,Bash,Glob,Grep (Opus)	—	DEFER	Overlaps prompt-engineer; only for real LLM pipelines.
embedded-systems	Read,Write,Edit,Bash,Glob,Grep	cross-agent list	DEFER	Only for hardware targets.
dependency-manager	Read,Write,Edit,Bash,Glob,Grep	wired context-mgr JSON + hard-coded vulnerabilities_fixed: 23 etc.	SKIP	Worst offender; structural rewrite needed.
ui-ux-tester	+chrome-mcp,computer-use	wired context-mgr JSON; hard runtime deps	SKIP	Team-orchestration design; deps you don't run.
build-engineer	Read,Write,Edit,Bash,Glob,Grep	context-mgr opener + fake % targets	SKIP	JS/monorepo framing ≠ UE5 build (UBT/CMake).
qa-expert	Read,Grep,Glob,Bash	context-mgr opener + baked coverage %	SKIP	Enterprise QA; overlaps code-review-worker + tdd-generator.
chaos-engineer	Read,Write,Edit,Bash,Glob,Grep	context-mgr opener	SKIP	No distributed-systems surface for a solo dev.
search-specialist	read-only	—	SKIP	Overlaps research-scout.
scientific-literature-researcher	+mcp__bgpt	BGPT MCP dependency	SKIP	Inert without the BGPT MCP.

All infrastructure (18 agents) and pure data-platform/ML agents: **SKIP** — no signal for a solo dev without production backend.

4. Recommended new workflows (recommendations only — not built)

Now that Dynamic Workflows are native, three are worth authoring as saved `/commands` when the moment comes. Cost amortizes across reruns, so build only when a workflow will be re-run.

1. **refresh-catalog** — promote the exact 4-source delta scan run for this report into a one-command pipeline: fan out one research agent per source → an adoption-verdict stage scoring each new agent against the kit's rules → a synthesized brief. Highest value; this report was assembled by hand and would be one command next time.

Pairs with the existing `refresh-cc-catalog` skill (which already forks a research-scout) — the workflow generalizes it to also produce verdicts.

2. **agent-adoption-bakeoff** — given a candidate upstream agent, fan out three lenses (harden-it / score-vs-kit-rules / check-overlap-with-existing) → synthesis verdict. Mirrors the proven `tdd-bakeoff.js` pattern already in GameMakerKit.
3. **portfolio-prelaunch-sweep** — chain the two DEFER reviewers (`accessibility-tester` + `penetration-tester`) as a gated pre-launch check when the portfolio ships. Read-only; one command, two perspectives.

Prompting discipline for all three: see §5 and `rules/dynamic-workflow-prompting.md` .

5. Better prompting for dynamic workflows (summary)

Full guide: `ClaudeKit/rules/dynamic-workflow-prompting.md` . The load-bearing points:

- **You can't steer mid-flight** — execution is synchronous. Control collapses to (1) the prompts up front, (2) between-phase barriers, (3) `/workflows` kill/resume. So the **prompt-review checkpoint** before scaling is the highest-ROI move.
- **Delegation anatomy** — every delegated task carries Objective · Output format · Tool/source guidance · Boundaries (scope *and* effort). Vague tasks are the #1 failure cause.
- **Pilot on 2–3, then scale.** A workflow wrong on item 1 is wrong 200x.
- **Bake the verify in** (a fresh agent refutes), and make the **schema demand evidence** (test output / diff / line ref), never a bare `{passed:true}` .
- **pipeline()** by default; **parallel()** only when a stage truly needs all prior results (dedup, early-exit, cross-comparison).
- **Upgrade the model before doubling the agent count** — tokens explain ~80% of perf variance; Fable 5 / Mythos 5 suit hard synthesis/verify stages.
- New craft folded in from shanraishan: **!command** live-state injection, **PROACTIVELY** triggers, **Agent(agent_type)** spawn restriction, and the **agent-teams data-contract** for mergeable parallel output.

6. Open questions & what wasn't verified

- **unslap maintenance / macOS-arm64 support** — unconfirmed; it's not currently installed, so `content-quality-editor` would be inert. Decision needed on whether AI-tell stripping belongs in the kit at all (vs. the deprecated `ai-writing-auditor`).
- **accessibility-tester line-95 context-manager block** — flagged but not read verbatim; confirm before any adoption (could flip DEFER → SKIP).
- **penetration-tester / qa-expert body tails** — only first ~50 lines read; SKIP/DEFER verdicts stand but the full bodies weren't audited.
- **shanraishan "Hot Features" dates & star counts** — not separable from automated commits; treat the list as aspirational until cross-checked against the official changelog.
- **project-artifact install path** — present in the `plugins/` dir but not confirmed in `marketplace.json` (so `/plugin install` may not yet resolve it), and it needs the Artifact beta regardless.
- **Unchecked VoltAgent depth** — `01-core-development` , `03-infrastructure` , `05-data-ai` raw files largely unread (dismissed as not-solo-relevant by category); a silent useful addition there is possible but unlikely.
- **CHARTER.md duplicate** — a second copy exists at `ClaudeCode/claude-kit/ProjectOptimizer/CHARTER.md` ; only the canonical `ClaudeCode/ProjectOptimizer/CHARTER.md` was updated this pass.